

# Lab 1 Report

Celyna Su (SID: 862037643) & Wei Xin Koh (SID: 862191103)

For Lab 1, we have implemented the following changes to the code:

1. Changed exit system call signature to void exit(int status) and implemented this change in user.h, defs.h, sysproc.h, prof.h, and other files that called exit()
  - a. We also changed the user space programs' call to this function from exit() to exit(0)

```
// Exit the current process. Does not return.
// An exited process remains in the zombie state
// until its parent calls wait() to find out it exited.
void
exit(int status)
{
    struct proc *curproc = myproc();
    struct proc *p;
    int fd;

    if(curproc == initproc)
        panic("init exiting");

    // Close all open files.
    for(fd = 0; fd < NOFILE; fd++){
        if(curproc->ofile[fd]){
            fileclose(curproc->ofile[fd]);
            curproc->ofile[fd] = 0;
        }
    }

    begin_op();
    iput(curproc->cwd);
    end_op();
    curproc->cwd = 0;

    acquire(&table.lock);

    // Parent might be sleeping in wait().
    wakeup1(curproc->parent);
    curproc->status = status;

    // Pass abandoned children to init.
    for(p = table.proc; p < &table.proc[NPROC]; p++){
        if(p->parent == curproc){
            p->parent = initproc;
            if(p->state == ZOMBIE)
                wakeup1(initproc);
        }
    }

    // Jump into the scheduler, never to return.
    curproc->state = ZOMBIE;
    sched();
    panic("zombie exit");
}
```

b.

2. Changed wait system call signature to `int wait(int* status)`
  - a. We deallocate the passed in status by setting `*status = p->status` in `sysproc.h`
  - b. We also changed the `wait()` function call to `wait(0)` in user space programs

```
// Wait for a child process to exit and return its pid.
// Return -1 if this process has no children.
int
wait(int *status)
{
    struct proc *p;
    int havekids, pid;
    struct proc *curproc = myproc();

    acquire(&ptable.lock);
    for(;;){
        // Scan through table looking for exited children.
        havekids = 0;
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->parent != curproc)
                continue;
            havekids = 1;
            if(p->state == ZOMBIE){
                // Found one.
                pid = p->pid;
                kfree(p->kstack);
                p->kstack = 0;
                freevm(p->pgdir);
                p->pid = 0;
                p->parent = 0;
                p->name[0] = 0;
                p->killed = 0;
                p->state = UNUSED;
                release(&ptable.lock);
                if(status != 0){
                    *status = p->status;
                }
            }
        }
        return pid;
    }
}

// No point waiting if we don't have any children.
if(!havekids || curproc->killed){
    release(&ptable.lock);
    return -1;
}

// Wait for children to exit. (See wakeup1 call in proc_exit.)
sleep(curproc, &ptable.lock); //DOC: wait-sleep
}
```

c.

3. Added a waitpid system call: int waitpid(int pid, int \*status, int options)
  - a. Added waitpid to SYSCALL to run it
  - b. Added waitpid to header files as well

```
int
waitpid(int pid, int *status, int options)
{
    struct proc *p;
    struct proc *curproc = myproc();
    int retrieve = 0;

    acquire(&ptable.lock);
    while(1)
    {
        retrieve = 0;
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
        {
            if(p->parent != curproc)
                continue;
            retrieve = 1;
            if(p->state == ZOMBIE && pid == p->pid && retrieve == 1){
                pid = p->pid;
                kfree(p->kstack);
                p->kstack = 0;
                freevm(p->pgdir);
                p->pid = 0;
                p->parent = 0;
                p->name[0] = 0;
                p->killed = 0;
                p->state = UNUSED;
                release(&ptable.lock);
                if(status != 0) {
                    *status = p->status;
                }
                return pid;
            }
        }
    }
    if(!retrieve || curproc->killed)
    {
        release(&ptable.lock);
        return -1;
    }
    sleep(curproc, &ptable.lock);
}
```

c.

4. Wrote test files for waitpid
  - a. Added a testbench titled lab1.c
  - b. Modified the Makefile by adding lab1.c to UPROGS and Extra

```
#include "types.h"
#include "user.h"

#define WNOHANG      1

int exitWait(void);
int waitPid(void);
int CELEBW02(void);

int main(int argc, char *argv[])
{
    printf(1, "\n This program tests the correctness of your lab#1\n");

    if (atoi(argv[1]) == 1)
        exitWait();
    else if (atoi(argv[1]) == 2)
        waitPid();
    else if (atoi(argv[1]) == 3)
        CELEBW02();
    else
        printf(1, "\ntype \"lab1 1\" to test exit and wait, \"lab1 2\" to test waitpid and \"lab1 3\" to test the extra credit WNOHANG option\n");

    // End of test
    exit(0);
    return 0;
}

int exitWait(void) {
    int pid, ret_pid, exit_status;
    int i;
    // use this part to test exit(int status) and wait(int* status)

    printf(1, "\n Parts a & b) testing exit(int status) and wait(int* status):\n");

    for (i = 0; i < 2; i++) {
        pid = fork();
        if (pid == 0) { // only the child executed this code
            if (i == 0)
            {
                printf(1, "\nThis is child with PID# %d and I will exit with status %d\n", getpid(), 0);
                exit(0);
            }
            else
            {
                printf(1, "\nThis is child with PID# %d and I will exit with status %d\n", getpid(), -1);
                exit(-1);
            }
        }
        else if (pid > 0) { // only the parent executes this code
            ret_pid = wait(&exit_status);
            printf(1, "\n This is the parent: child with PID# %d has exited with status %d\n", ret_pid, exit_status);
        } else // something went wrong with fork system call
        {
            printf(2, "\nError using fork\n");
            exit(-1);
        }
    }
    return 0;
}
```